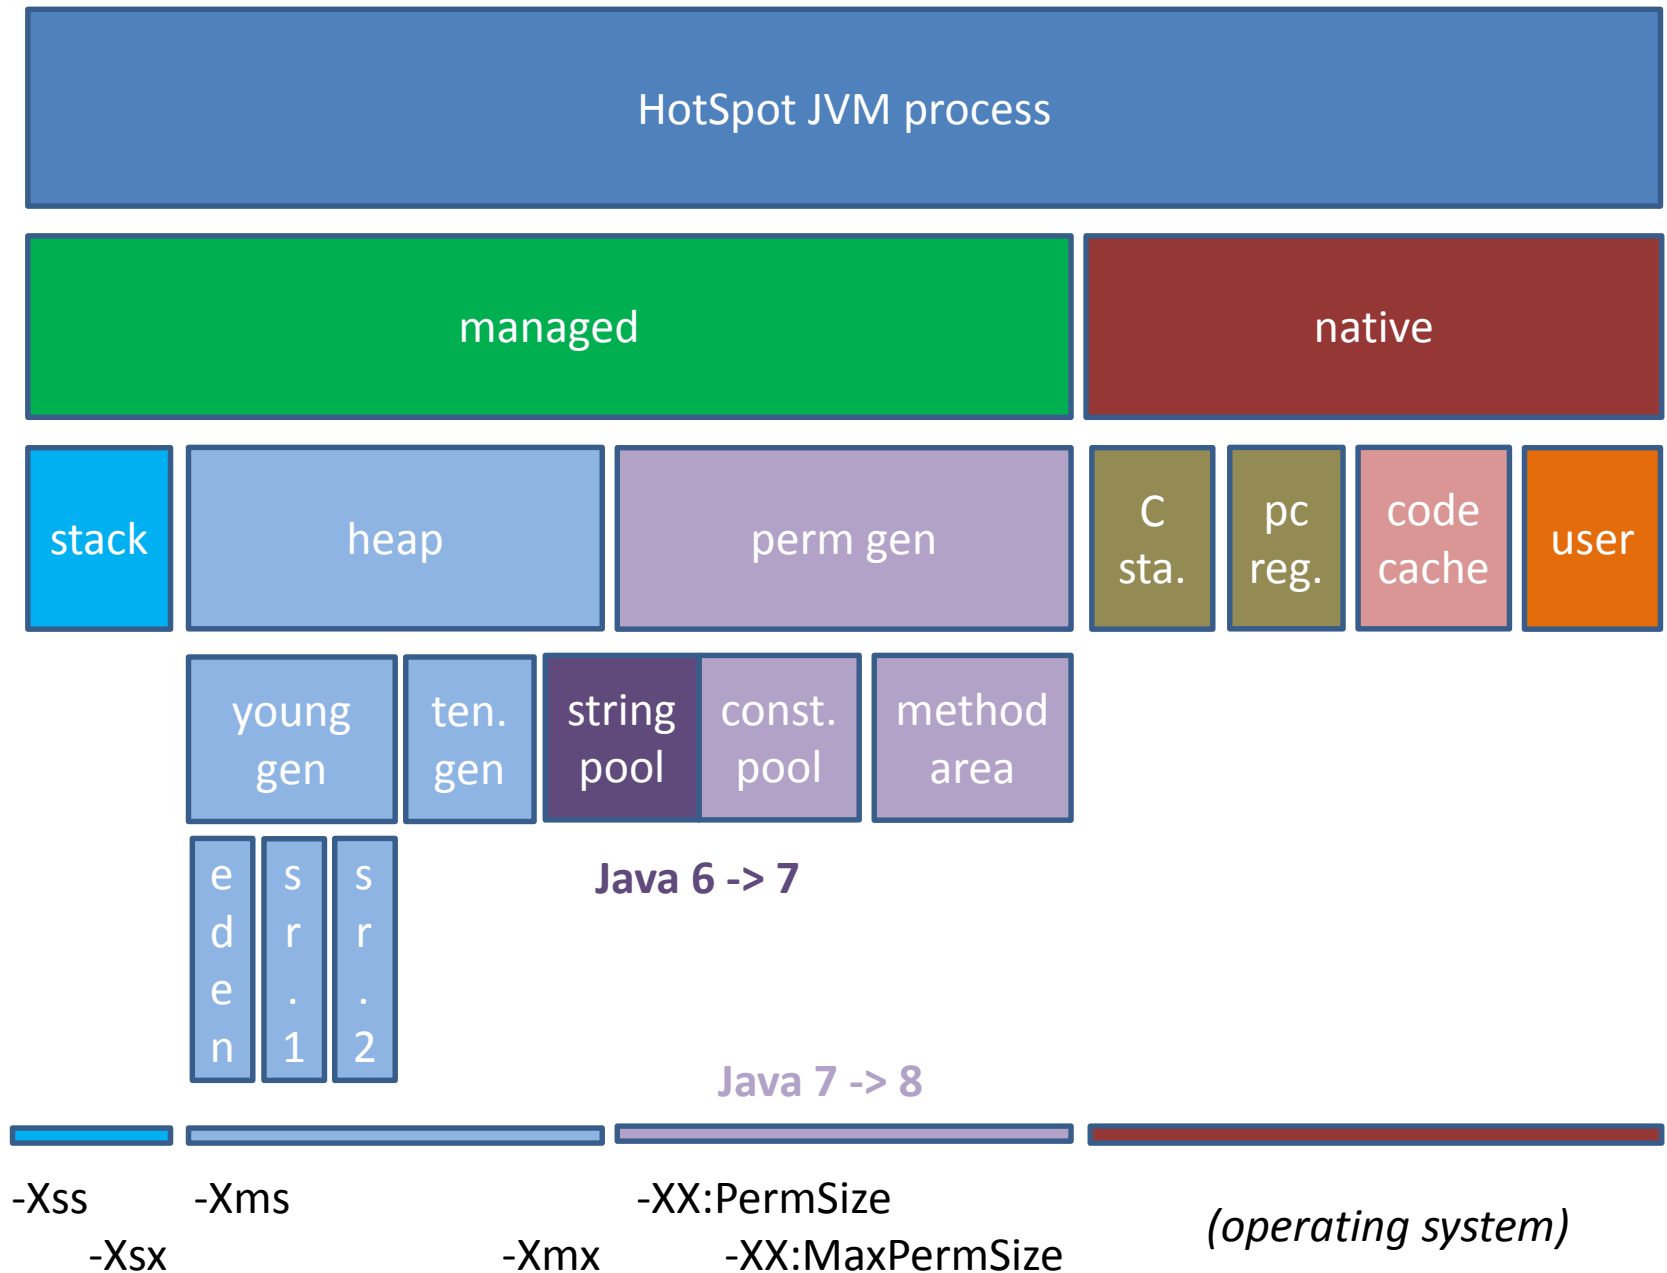
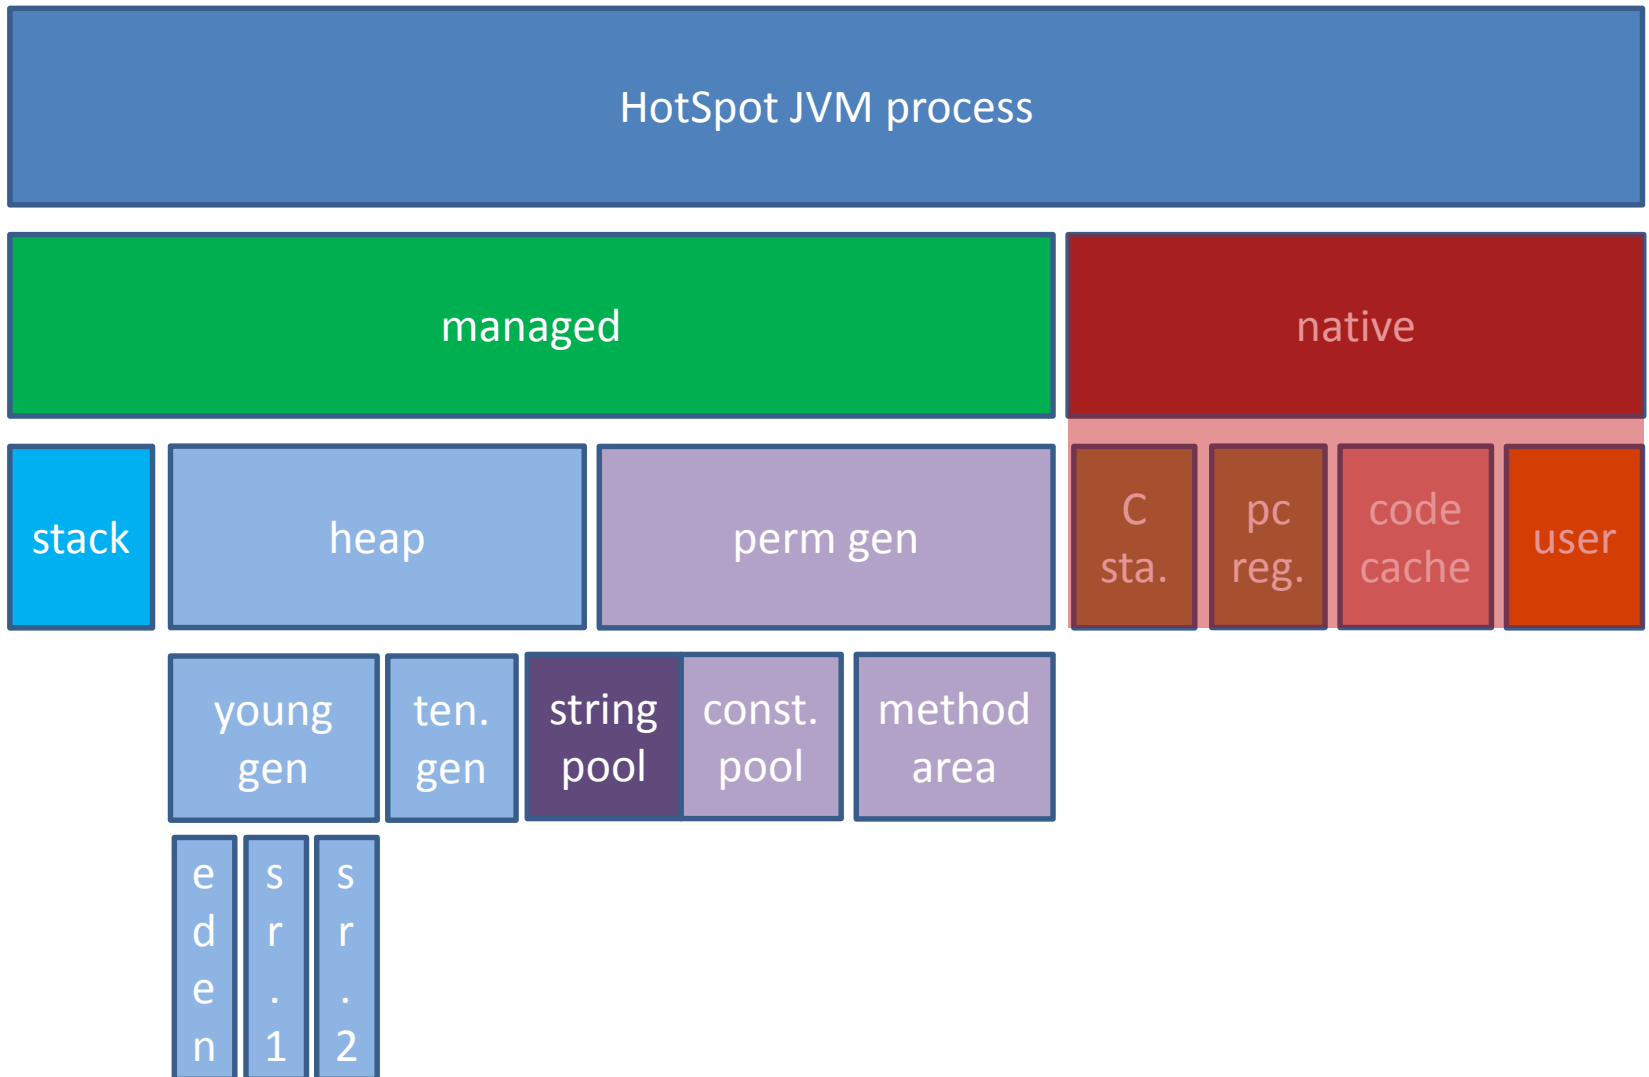


A topology of memory leaks on the JVM



genuine leaks



Allocating native memory with *sun.misc.Unsafe*

1. Getting hold of “the unsafe”

```
Field unsafe = Unsafe.class
    .getDeclaredField("theUnsafe");
unsafe.setAccessible(true);
Unsafe theUnsafe = (Unsafe) unsafe.get(null);
```

2. Allocating native memory

```
final int intSize = 4;
long arraySize = (1L + Integer.MAX_VALUE) * intSize;
long index = unsafe.allocateMemory(arraySize);
Random random = new Random();
for(long l = 0L; l < arraySize; l++)
    unsafe.putInt(random.nextInt());
```

3. (Hopefully) releasing native memory

```
// Without me, the memory leaks
unsafe.freeMemory(index);
```

I don't do *sun.misc.Unsafe*

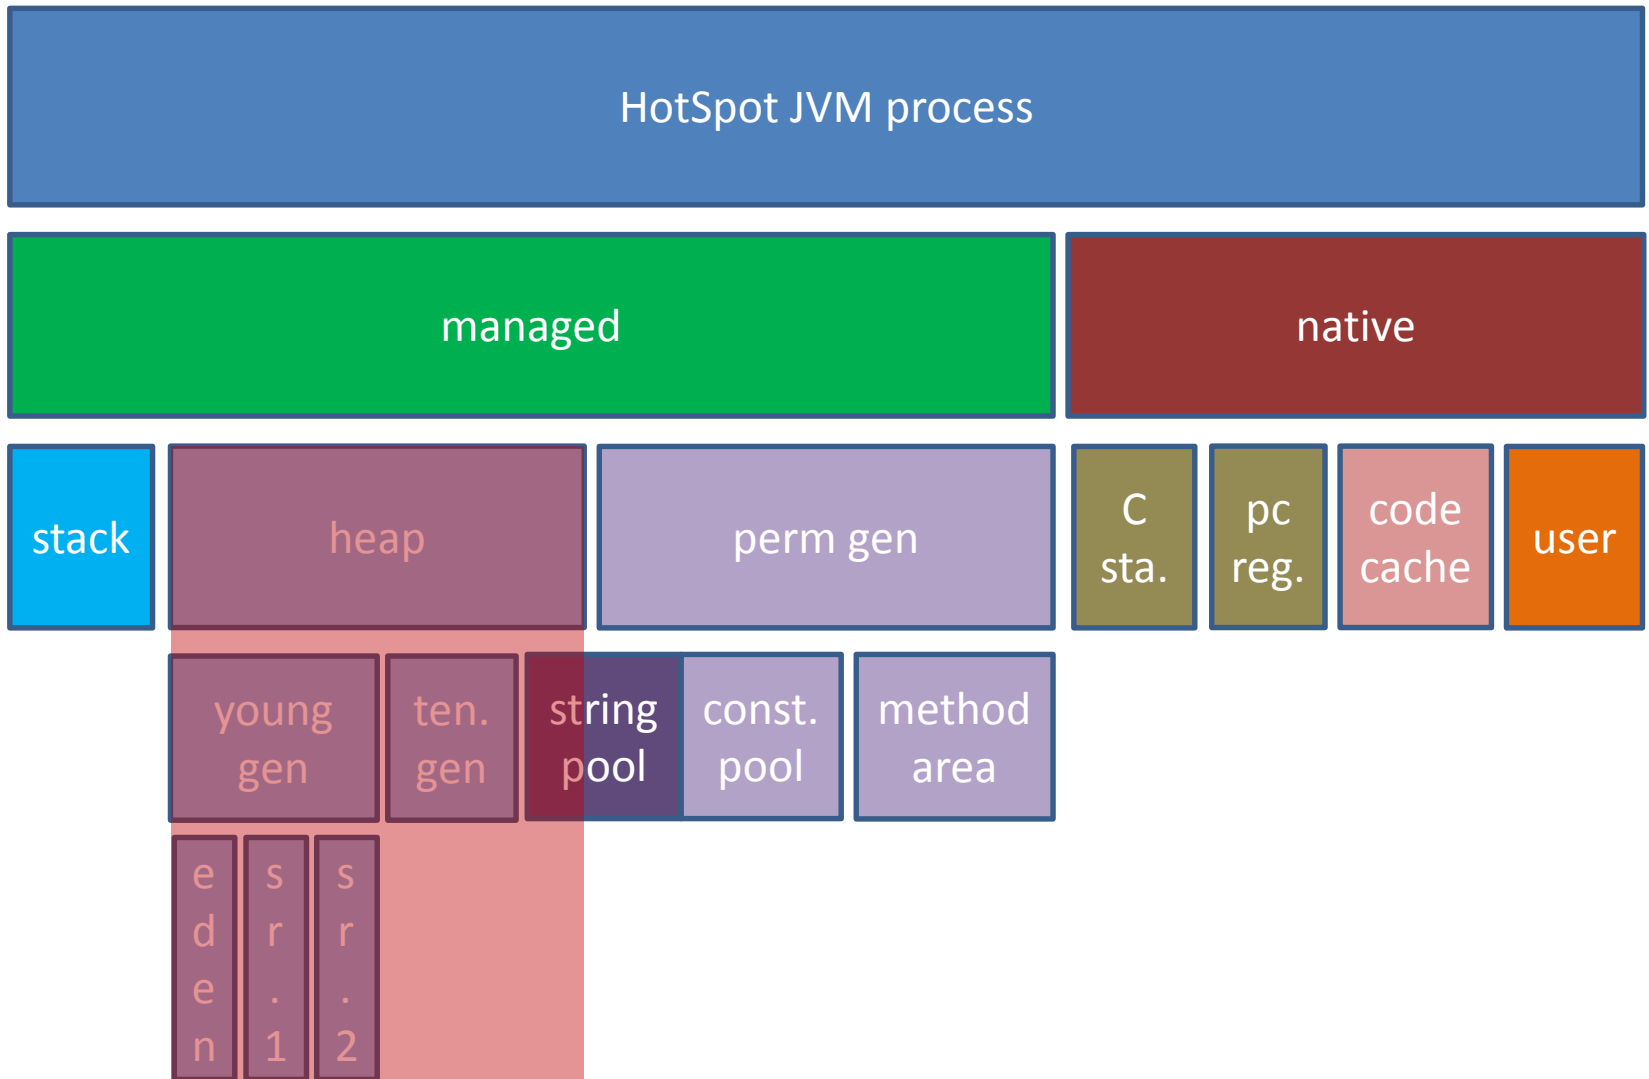
1. But some of your favorite tools do:
 1. Kryo (used by e.g. Twitter's Storm, Apache Hive)
 2. EhCache's "big memory" (used by e.g. Hibernate)
 3. JDK (e.g. direct byte buffers)
 4. General purpose (e.g. GSON, Snappy)
2. *Unsafe* to become public API in Java 9 (competition with CLR?)
3. JNI leaks have the same effect
4. Java 8 removes "perm gen" in favor of native memory "meta space"

symptoms of a genuine leak

No exception required by JVM. However, JVM will be untypical slow when:

1. Creating a thread:
Requires allocation of new stacks / pc register.
2. Loading a class / JIT:
Requires native memory.
3. Doing I/O:
Often requires native memory.
4. Calling a native method:
Allocates resources on native stack / user space. Swapping can delay execution.
5. Garbage collection:
Tenured generation collection needs broad access such that OS must swap back the heap. Rule of thumb: All JVM heaps must be swapped back into memory. (Leaked memory will stay swapped out.)

Heap leaks



```
class FixedSizeStack {
```


```
    private final Object[] objectPool = new Object[10];  
    private int pointer = -1;
```

```
    public Object pop() {  
        if(pointer < 0)  
            throw new NoSuchElementException();  
        return objectPool[pointer--];  
    }
```

```
    public Object peek() {  
        if(pointer < 0)  
            throw new NoSuchElementException();  
        return objectPool[pointer];  
    }
```


```
    public void push(Object object) {  
        if(pointer > 8)  
            throw new IllegalStateException("stack overflow");  
        objectPool[++pointer] = object;  
    }  
}
```





All memory leaks in Java are linked to **reference** life cycles.

**always prefer
immutability**




JDK 6

```
public String substring(int beginIndex, int endIndex) {  
    // omitted argument validation  
    return ((beginIndex == 0) && (endIndex == count))  
        ? this  
        : new String(offset + beginIndex,  
                      endIndex - beginIndex,  
                      value);  
}
```

```
String(int offset, int count, char[] value) {  
    this.value = value;  
    this.offset = offset;  
    this.count = count;  
}
```

XML, anybody?

```
org.xml.sax.DocumentHandler {  
    // several callback methods  
    void characters(char[] ch, int start, int length);  
}
```



Non-static inner classes

```
new Message() {  
    @Override public String getInfo() {  
        return "bar";  
    }  
}
```

avoid non-static
inner classes

uncompiled

desugared

```
class ExampleActivity$1 implements Message {  
  
    private ExampleActivity this$0;  
  
    ExampleActivity$1(ExampleActivity this$0) {  
        this.this$0 = this$0;  
    }  
  
    @Override public String getInfo() {  
        return "bar";  
    }  
}
```

Java 8 lambda expressions

```
class Foo {  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        list.forEach(s -> { System.out.println(s); });  
    }  
}
```

uncompiled

```
class Foo {  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        list.forEach(LambdaMetafactory  
            .INVOKEDYNAMIC(Foo::lambda$1)  
            .make());  
    }  
}
```

desugared

```
private static void lambda$1(String s) {  
    System.out.println(s);  
}  
}
```


Java 8 lambda expressions

```
class Foo {  
    final String prefix; // constructor omitted  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        list.forEach(s -> { System.out.println(  
            prefix + ":" + s) });  
    }  
}
```

uncompiled

```
class Foo {  
    final String prefix; // constructor omitted  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        list.forEach(LambdaMetaFactory  
            .INVOKEDYNAMIC(this::lambda$1)  
            .make());  
    }  
}
```

desugared

```
private void lambda$1(String s) {  
    System.out.println(this.prefix + ":" + s);  
}
```

Java 8 lambda expressions

```
class Foo {  
    final String prefix; // constructor omitted  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        String prefix = this.prefix;  
        list.forEach(s -> { System.out.println(  
            prefix + ":" + s) });  
    }  
}
```

pay attention to
your lambda
expression's scope

uncompiled

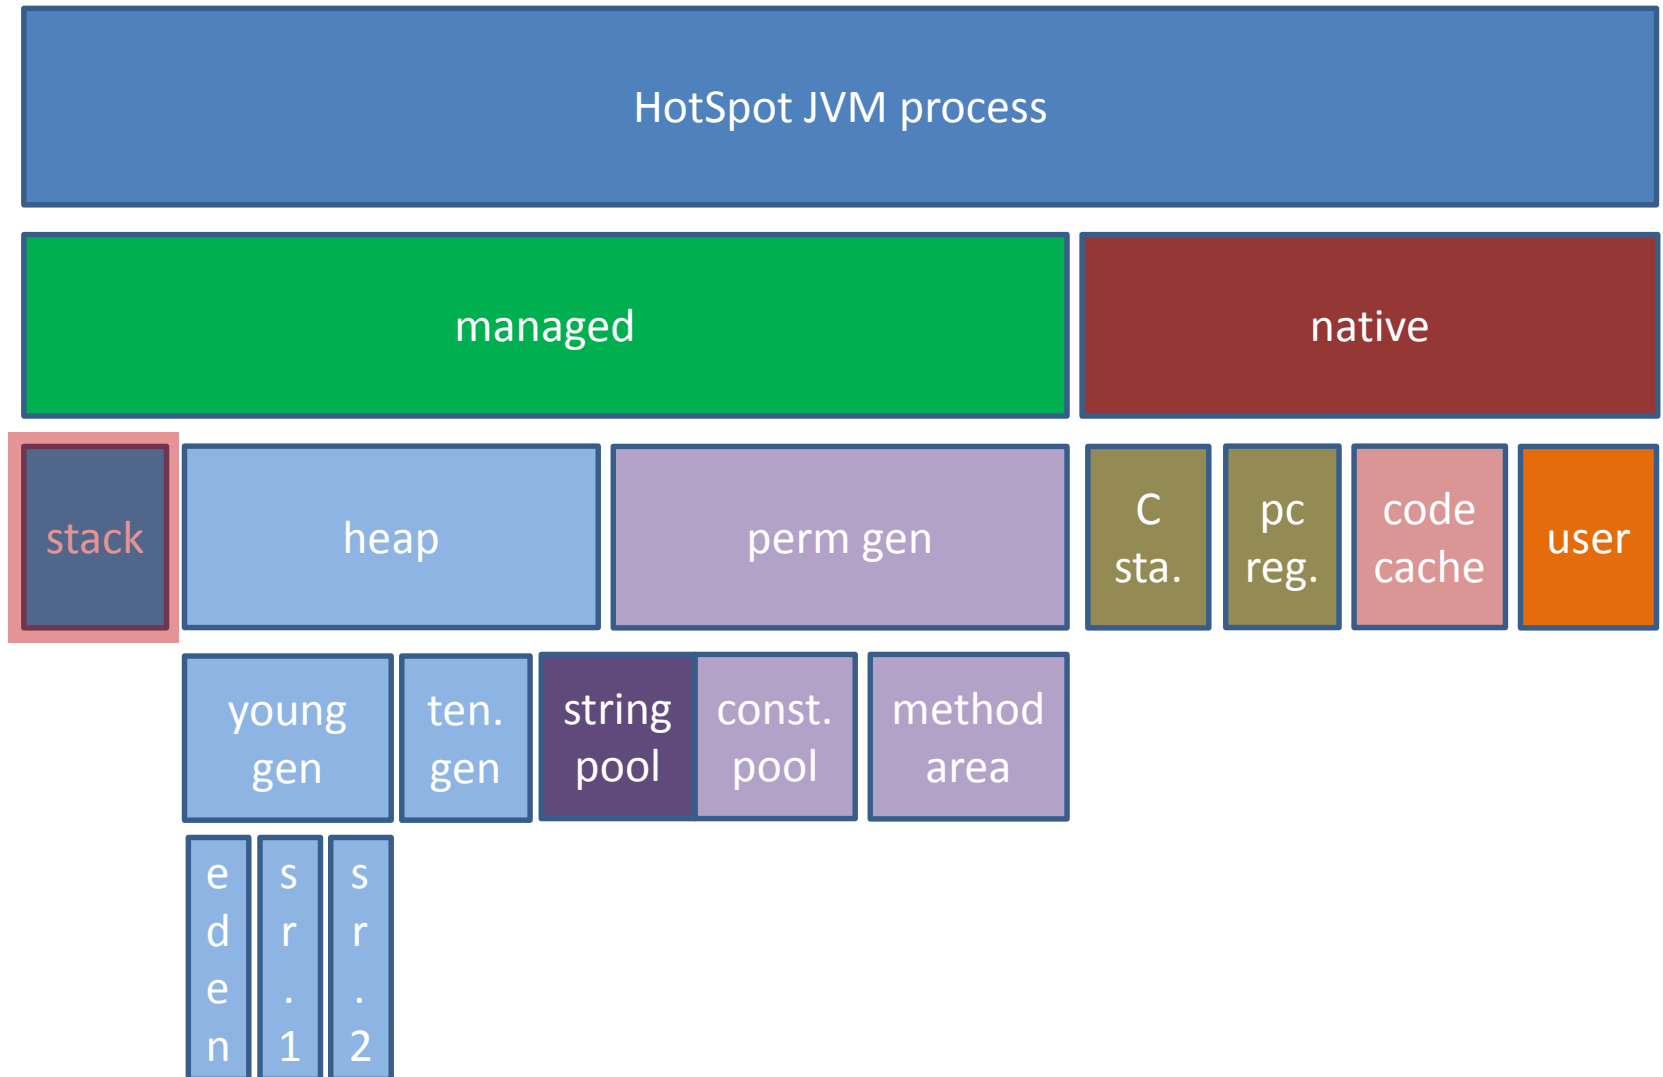
```
class Foo {  
    final String prefix; // constructor omitted  
    void bar() {  
        List<String> list = Arrays.asList("foo", "bar");  
        String prefix = this.prefix;  
        list.forEach(LambdaMetafactory  
            .INVOKEDYNAMIC(Foo::lambda$1)  
            .make(prefix));  
    }  
    private static void lambda$1(String prefix, String s) {  
        System.out.println(prefix + ":" + s);  
    }  
}
```

desugared

Other typical causes of heap leaks

- Functional expressions in e.g. Scala / Groovy
- Serialization leaks (e.g. Apache Wicket)
- Singletons / enums
- Context frameworks (e.g. DI like Spring)

Stack leaks






```
class StackLeak {
    int SIZE = (int) (0.5 * Runtime
                      .getRuntime()
                      .maxMemory());

    void foo() {
        {
            byte[] array1 = new byte[SIZE];
        }
        byte[] array2 = new byte[SIZE];
    }


    void bar() {
        {
            byte[] array1 = new byte[SIZE];
        }
        byte[] array2 = new byte[10];
        byte[] array3 = new byte[SIZE];
    }
}
```



```
void foo() {  
    {  
        byte[] array1 = new byte[SIZE];  
    }  
    byte[] array2 = new byte[SIZE];  
}
```

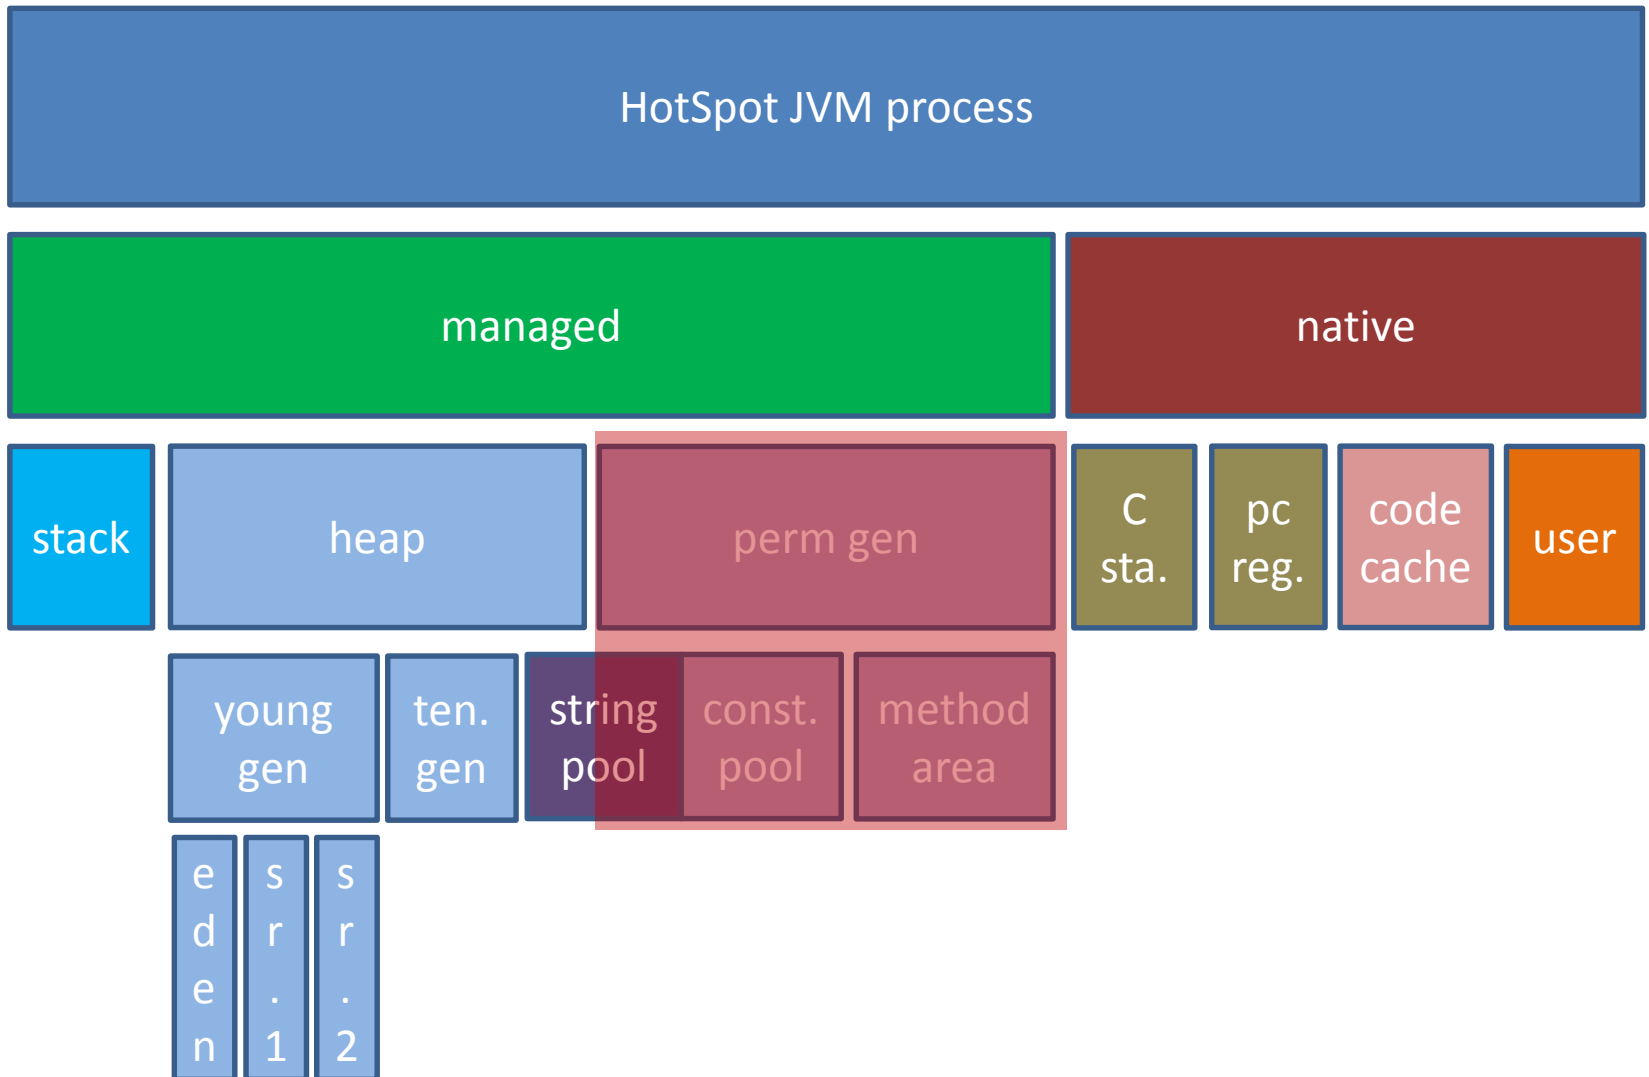
Two black arrows point to the first and second array declarations in the code block above.

write short methods,
mistrust explicit scopes

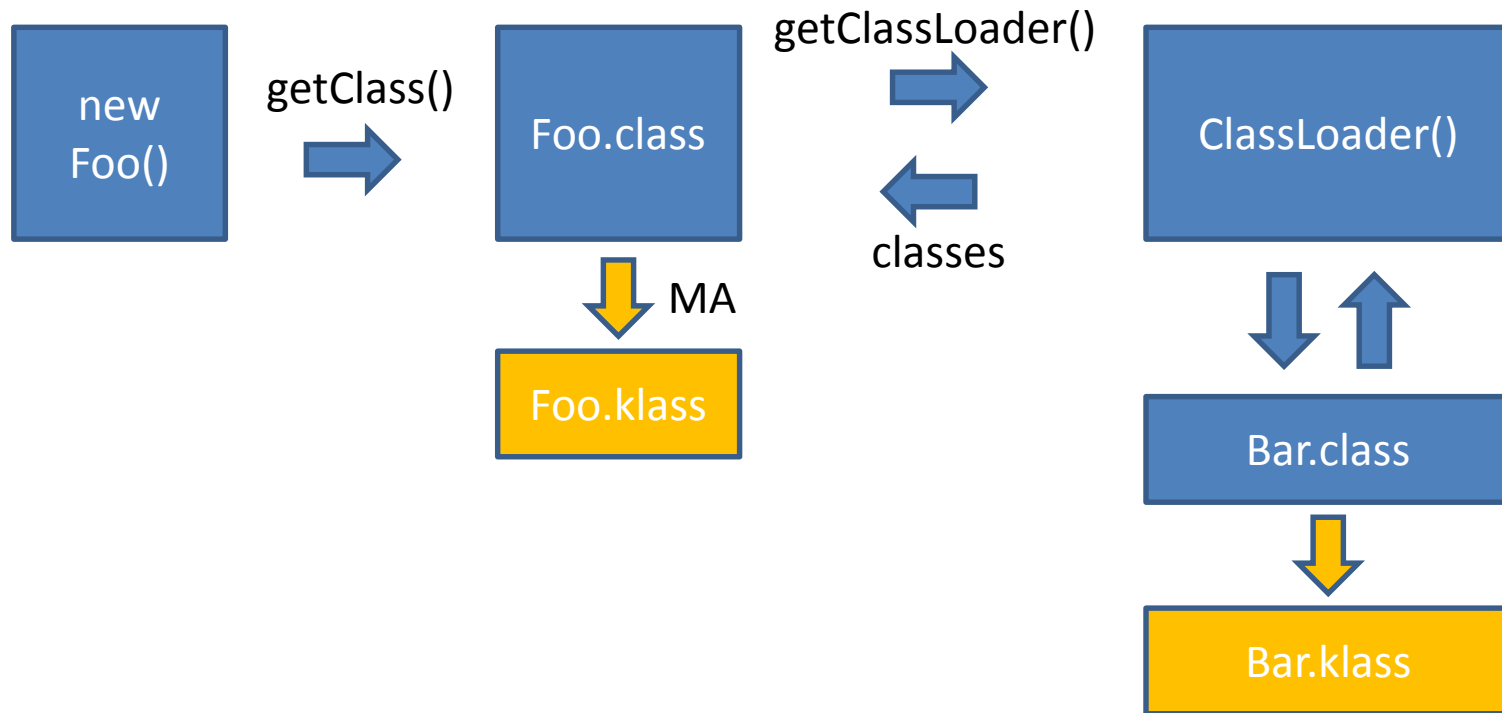


```
void bar() {  
    {  
        byte[] array1 = new byte[SIZE];  
    }  
    byte[] array2 = new byte[1];  
    byte[] array3 = new byte[SIZE];  
}
```

“Perm gen” leaks

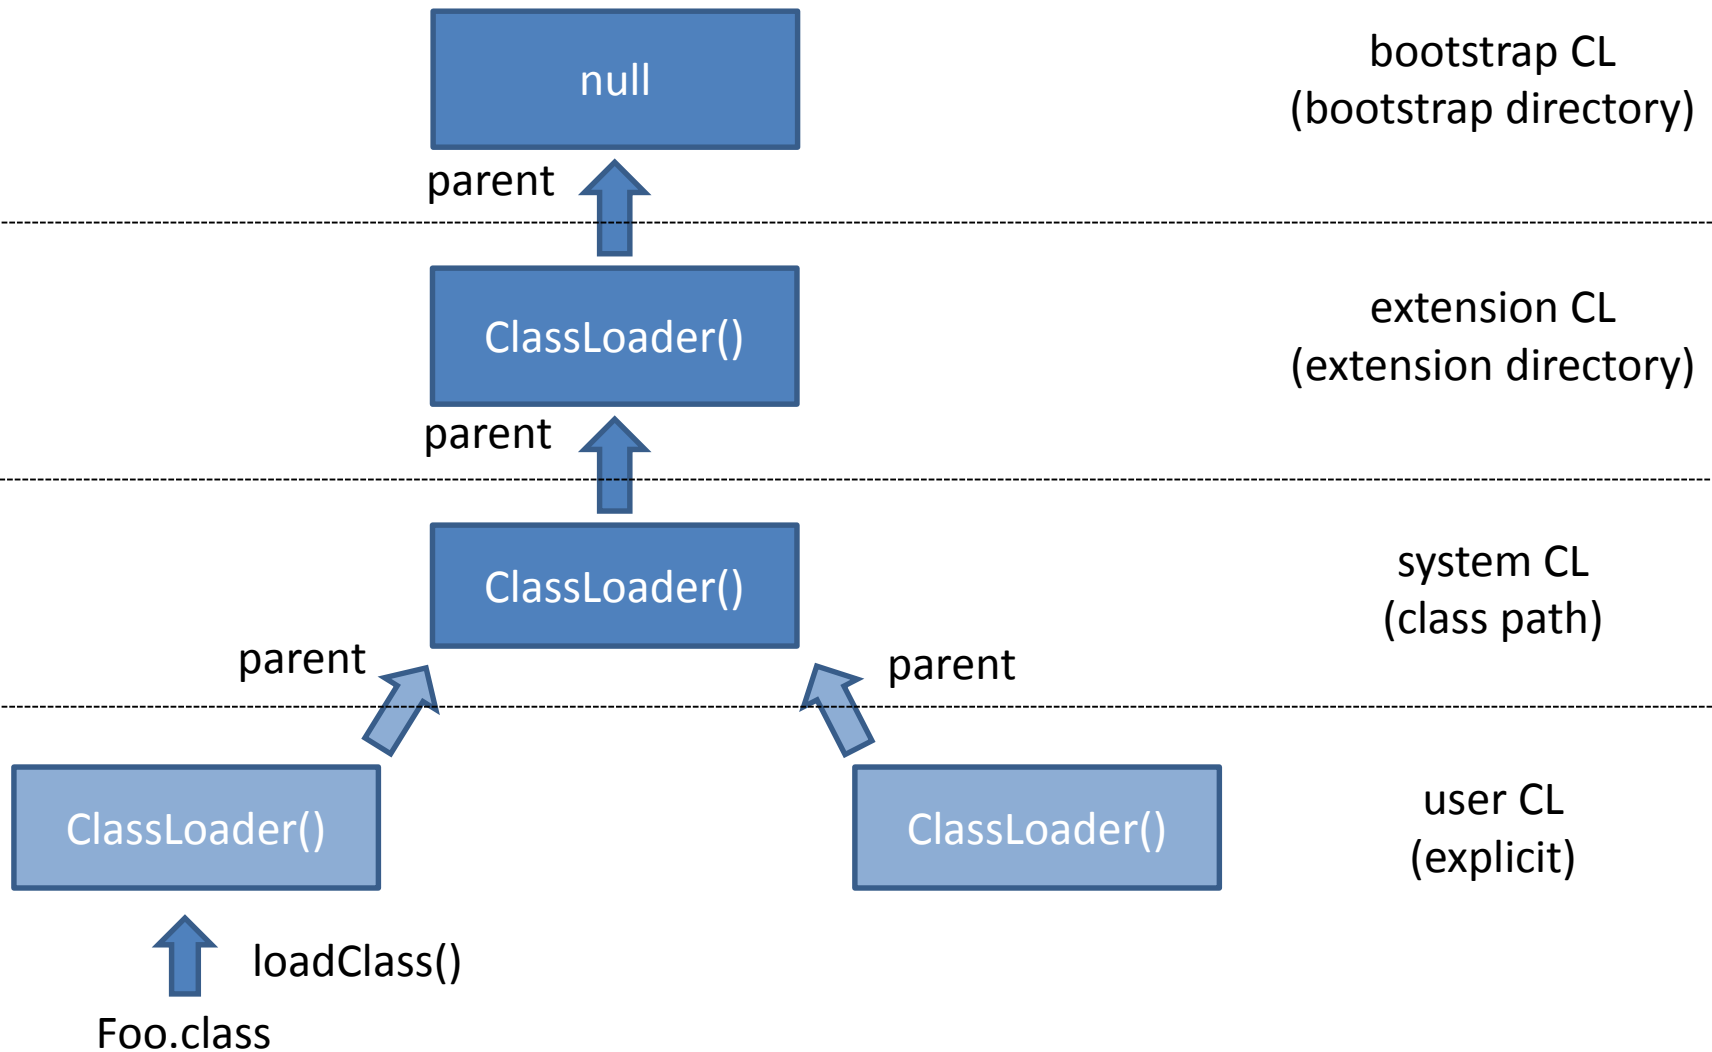


Classes and HotSpot

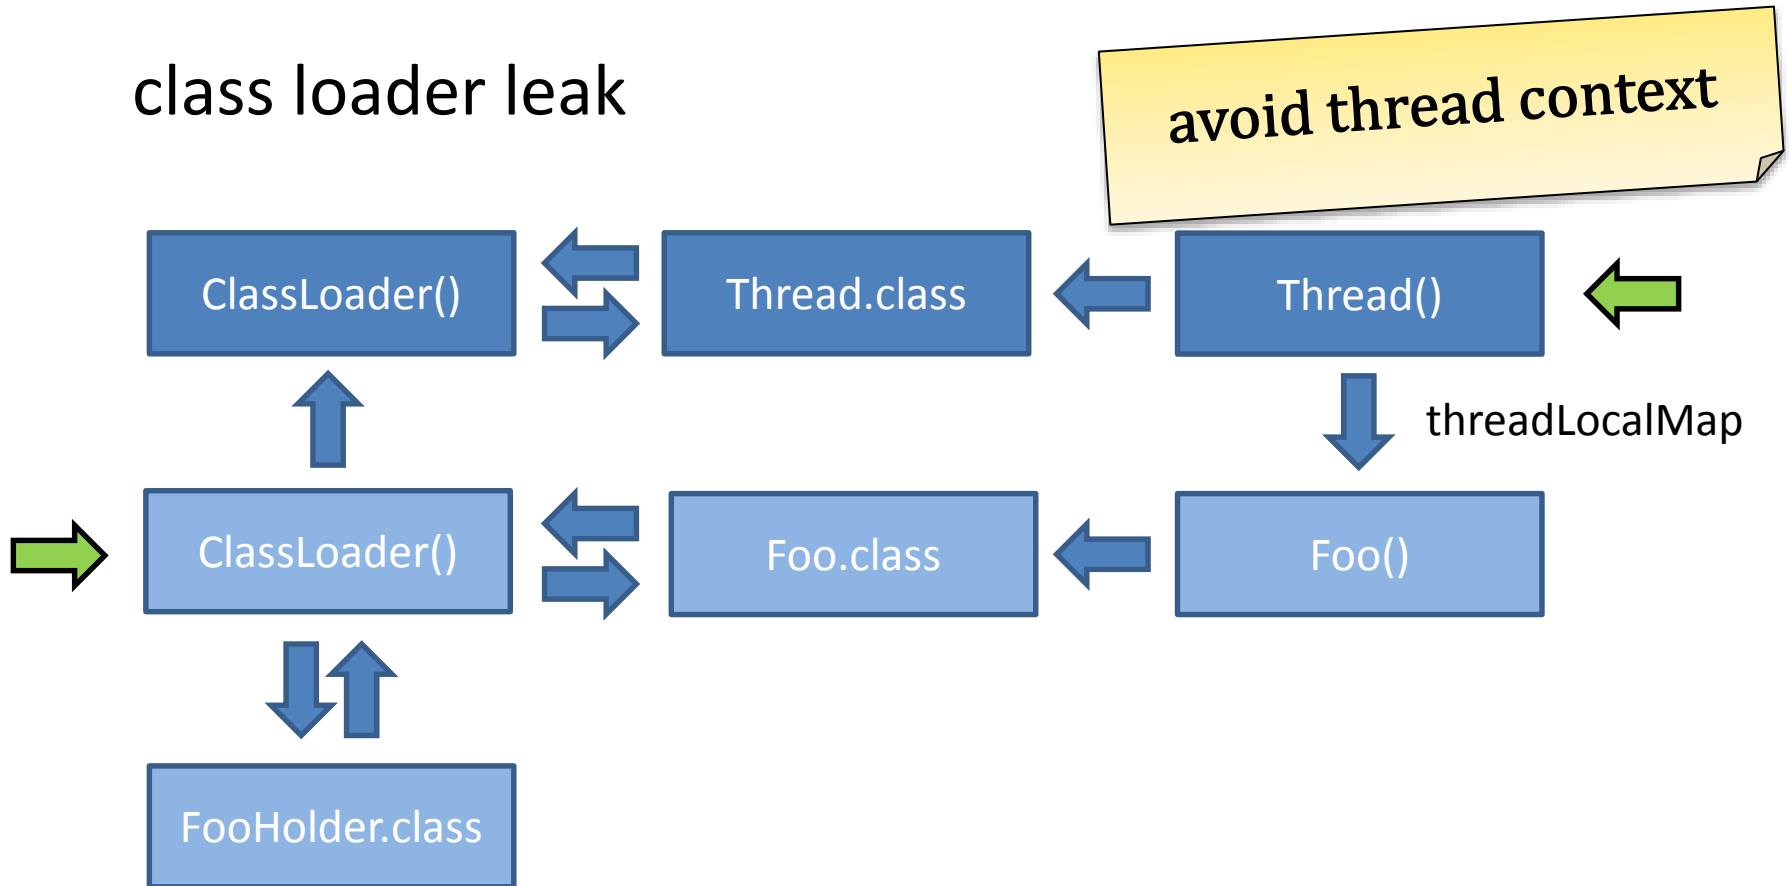


```
Field field = ClassLoader.class.getDeclaredField("classes");
field.setAccessible(true);
Vector<Class<?>> classes = (Vector<Class<?>>)
    field.get(ClassLoader.getSystemClassLoader());
```

ClassLoaders and HotSpot



class loader leak



```
class FooHolder {  
    static ThreadLocal<Foo> tlFoo = new ThreadLocal<Foo>();  
    static { tlFoo.set(new Foo()); }  
}
```

Other typical causes of

avoid redeployment,
use one container per app

- Shut down hooks
- Use of thread context class loaders (e.g. OSGi)
- Service provider interfaces (SPIs)
- JDBC drivers (JDK *DriverManager*)
- Security frameworks (e.g. JDK *Policy*)
- Class loader magic (e.g. instrumentation)



Why do modern applications require so much perm gen memory?

```
class Foo {  
    public Object bar(Object o) { return o; }  
}
```

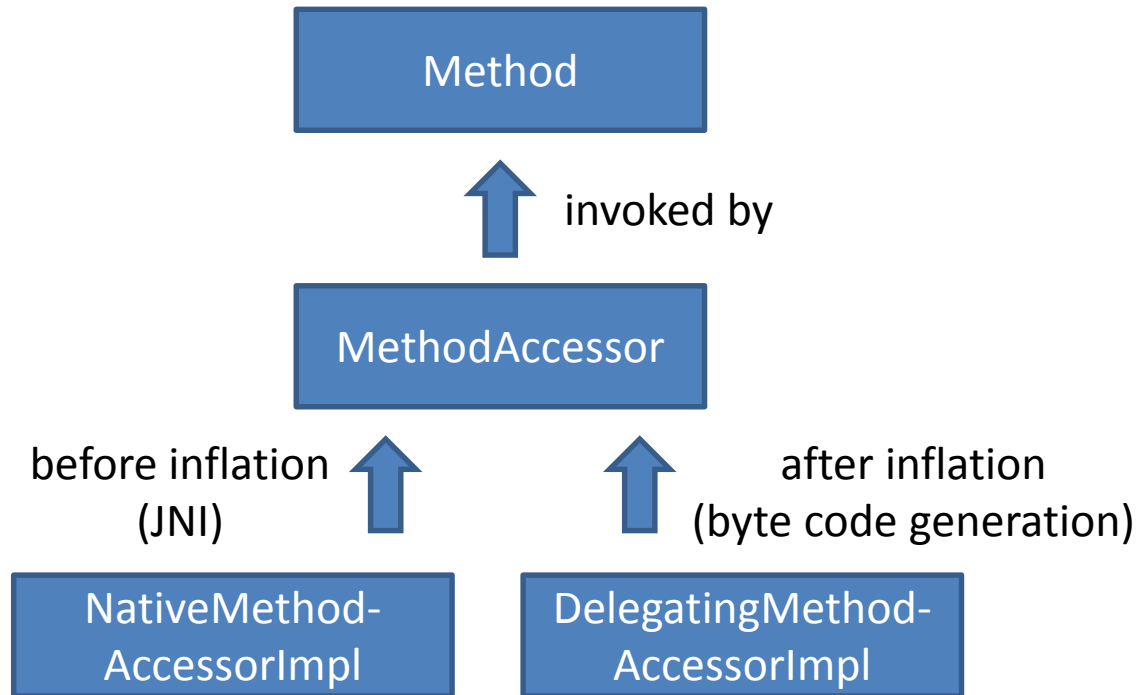
```
Enhancer enhancer = new Enhancer();  
enhancer.setSuperclass(Foo.class);  
enhancer.setCallback(new MethodInterceptor() {  
    @Override public Object intercept(Object obj,  
        Method method,  
        Object[] args,  
        MethodProxy proxy) throws Throwable {  
        return proxy.invokeSuper(obj, doSomethingWith(args));  
    }  
});  
Foo enhancedFoo = (Foo) enhancer.create();
```

Created classes: $2 + \text{\#methods} * 2$
[FastClass, FastMethod, MethodProxy]

avoid instrumentation

Modern JDK's inflation works similarly

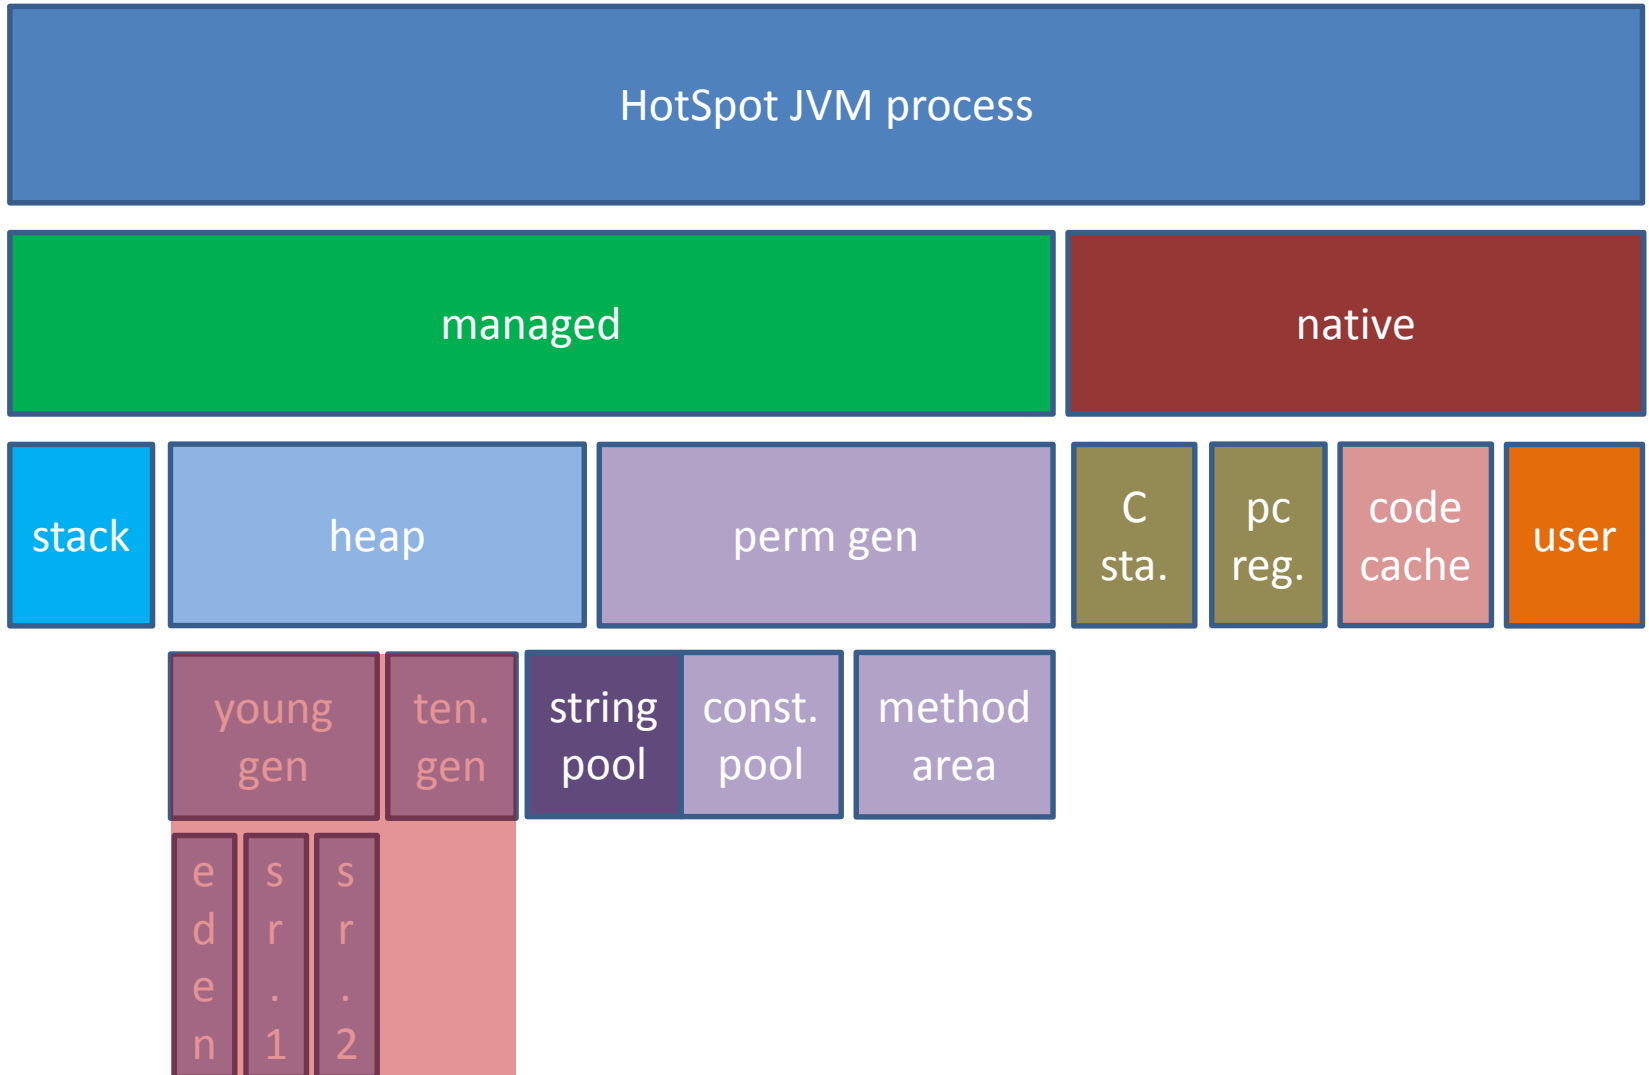
remember inflation



Java 8 meta space

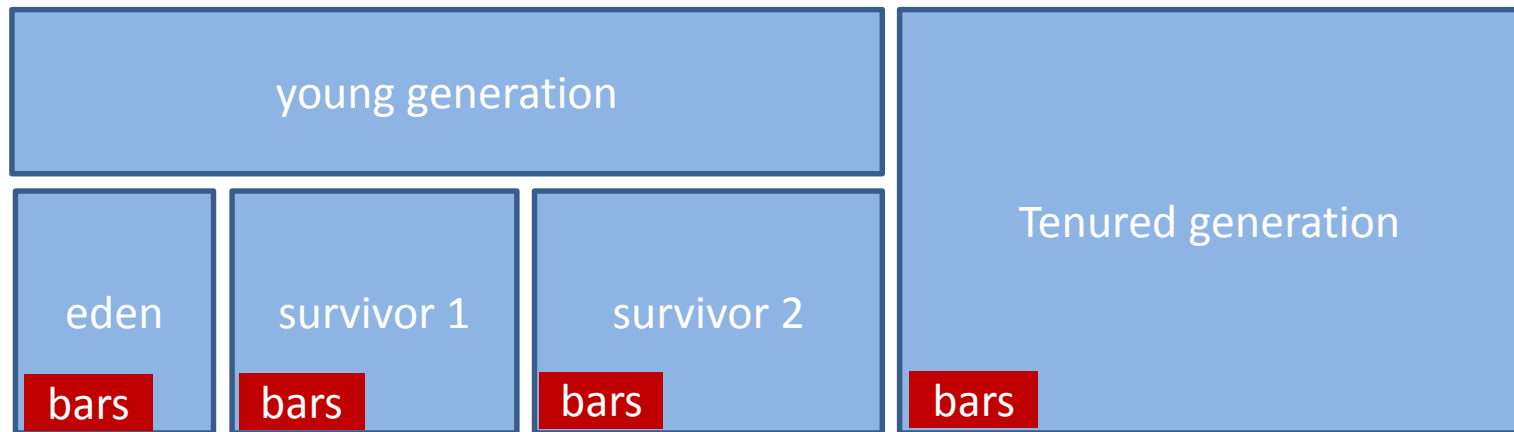
- Permanent generation rebranded (and probably approximation to JRockit)
- Meta space is allocated in native memory
- No space limit by default (*MaxMetaspaceSize*)
- Today's debugging tools will not be able to read meta space

Tenuring leaks



don't be scared of "new",
avoid object pooling,
JIT knows best

```
public void foo() {  
    for(int i = 0; i < 100; i++) {  
        Set<Bar> bars = new HashSet<Bar>();  
        for(int j = 0; j < 100; j++) {  
            bars.add(makeBar(i, j));  
        }  
        doSomethingWith(bars);  
    }  
}
```



collection of survivor 2

Implicit memory leaks

- Leaked **threads**: Require fixed amount of memory for call stacks / pc register / general allocation
- Leaked **handles** (files, sockets, databases): Usually require JVM memory resources

debugging / profiling

```
jmap -dump:live,format=b,file=<...> <pid>
```

```
java -XX:+HeapDumpOnOutOfMemoryError  
-XX:+HeapDumpPath=<...> <...>
```

```
jhat <hprof file>  
select a from [I a where a.length >= 256
```

GUI tools: MAT (Eclipse RCP) / JVisualVM / HeapWalker

**tools that create heap dumps
by instrumentation require
memory to run**

<http://rafael.codes>
[@rafaelcodes](#)



<http://documents4j.com>
<https://github.com/documents4j/documents4j>



<http://bytebuddy.net>
<https://github.com/raphw/byte-buddy>

